

Li Zhang

+1 412-328-6886 | liz239@pitt.edu | [LinkedIn](#) | [Google Scholar](#)

EDUCATION

University of Pittsburgh

Pittsburgh, PA, U.S.A.

Doctor of Philosophy in Intelligent Systems (Artificial Intelligence)

Aug. 2024 - Apr. 2028

- Research: NLP for legal reasoning, trustworthy LLMs, agentic AI, and system evaluation
- Advisor: Prof. [Kevin Ashley](#)
- Courses: Large Language Models (Carnegie Mellon University), Natural Language Processing

Tsinghua University

Beijing, China

Master of Computational Law

Sep. 2019 - Jun. 2022

- Research: Computational Law

Australian National University

Canberra, Australia

Study Abroad in College of Engineering and Computer Science

Jul. 2017 - Dec. 2017

- GPA: 6.5/7.0

China University of Mining and Technology

Jiangsu, China

Bachelor of Electrical Engineering

Sep. 2015 - Jun. 2019

- GPA: 90.91/100, Ranking: 2/278

TECHNICAL EXPERIENCE

Thomson Reuters Labs

Eagan, MN, U.S.A.

Applied Research Intern

May 2026 - Aug. 2026

- Conducted applied research on long-term memory for AI agents, supporting deep research over long legal documents and long-form content generation with direct commercial applications in litigation and transactional workflows.
- Built and evaluated agent pipelines across dataset optimization, model adaptation, and automated assessment; partnered with science, product, and engineering teams to deliver tailored solutions for grounded workflows.

SELECTED PUBLICATIONS

1. DLawBench: Evaluating LLMs Through Multi-Turn Legal Consultation [\[Link\]](#)

Li Zhang, Yuzhen Shi, Yiran Hu, Jingwen Zhang, Wenbo Lv, Yubo Ma, Wei Wang, Rongyao Shi, Yuanyang Qiu, Xinran Xu, Yuemeng Qi, Linlin Miao, Jaromir Savelka, Yun Liu, Kevin Ashley, Bing Zhao, Hu Wei, Lin Qu
Preprint, under review

2. Thinking Longer, Not Always Smarter: Evaluating LLM Capabilities in Hierarchical Legal Reasoning [\[Link\]](#)

Li Zhang, Matthias Grabmair, Morgan Gray, Kevin Ashley
CS&Law 2026

3. PLawBench: A Rubric-Based Benchmark for Evaluating LLMs in Real-World Legal Practice [\[Link\]](#)

Yuzhen Shi, Huanghai Liu, Yiran Hu, Song gaojie, Xu Xinran, Yubo Ma, Tianyi Tang, Li Zhang, Qingjing Chen, Feng Di, Wenbo Lv, Weiheng Wu, Kexin Yang, Sen Yang, Wei Wang, Rongyao Shi, Qiuyuan Yang, Yuemeng Qi, Jingwen Zhang, Xiaoyu Sui, Yifan Chen, Yi Zhang, An Yang, Bowen Yu, Dayiheng Liu, Junyang Lin, Weixin Shen, Bing Zhao, Charles Clarke, Hu Wei
ACL 2026

4. Retrieval-Based Multi-Label Legal Annotation: Extensible, Data-Efficient and Hallucination-Free [\[Link\]](#)

Li Zhang, Jaromir Savelka, Kevin Ashley
ICAIL 2026

5. DenseMixer: Improving MoE Post-Training with Precise Router Gradient [\[Link\]](#)

Feng Yao, Junxia Cui, Ruohan Zhang, Liyuan Liu, Shibo Hao, Li Zhang, Chengyu Dong, Shuohang Wang, Yelong Shen, Jianfeng Gao, Jingbo Shang
NeurIPS 2025 Workshop on Structured Probabilistic Inference & Generative Modeling

6. Do LLMs Truly Understand When a Precedent Is Overruled? [\[Link\]](#)

Li Zhang, Jaromir Savelka, Kevin Ashley
JURIX 2025, Best Student Paper

7. **Mitigating Manipulation and Enhancing Persuasion: A Reflective Multi-Agent Approach for Legal Argument Generation** [\[Link\]](#)
Li Zhang, Kevin Ashley
ICAIL 2025 Workshop
8. **Generating Case-Based Legal Arguments with LLMs** [\[Link\]](#)
Morgan Gray, **Li Zhang**, Kevin Ashley
CS&Law 2025
9. **Intelligent Proximate Analysis of Coal Based on Near-Infrared Spectroscopy** [\[Link\]](#)
Weinan Liu, Boyuan Peng, Xiaoyu Liu, Feifan Ren, **Li Zhang** (corresponding author), *Journal of Applied Spectroscopy*

RESEARCH PROJECTS

- Benchmarking LLM Long-Context Understanding and Temporal Logic** Pittsburgh, PA, U.S.A.
Li Zhang, Jaromir Savelka, Kevin Ashley Sep. 2024 – Present
- Constructed a high-stakes benchmark from real-world legal corpora to evaluate long-context understanding and temporal reasoning over complex overruling relationships – a previously unsolved evaluation problem.
 - Designed quantitative metrics for temporal bias, robustness, and hallucination, ensuring generated case histories remain grounded in the controlling precedent of record.
- Reflective Multi-Agent Frameworks for Ethical Legal Persuasion** Pittsburgh, PA, U.S.A.
Li Zhang, Kevin Ashley Aug. 2024 – Present
- Architected custom multi-agent algorithms with reflective verification, where independent critic agents test generated arguments against factual records to minimize hallucination.
 - Implemented inspectable guardrails ensuring systems advocate effectively without manipulating facts or fabricating evidence – controls designed for regulated, high-stakes settings.
- Evaluating and Enhancing Hierarchical Legal Reasoning** Pittsburgh, PA, U.S.A.
Li Zhang, Matthias Grabmair, Morgan Gray, Kevin Ashley Nov. 2024 – Present
- Investigated reasoning collapse in LLMs through systematic evaluation of multi-stage reasoning tasks requiring case distinctions based on subtle factual predicates.
 - Developed hierarchical inference strategies that improve the robustness of nuanced logical deduction in complex case-law scenarios.

FINANCIAL & LEGAL EXPERIENCE

- The Export-Import Bank of China, Jiangsu Branch** Nanjing, China
Legal Counselor Jun. 2022 - Jul. 2024
- Led negotiation and execution of cross-border structured financing and trade-finance transactions at a state policy bank, advising on credit risk, counterparty exposure, and regulatory compliance.
 - Drafted and reviewed credit facility and security documentation; advised deal teams on transaction structuring, operational risk, and dispute resolution, building first-hand knowledge of financial markets and banking products.

TALKS

1. Do LLMs Truly Understand When a Precedent Is Overruled, [2025 Pittsburgh Legal AI Summit](#), Carnegie Mellon University
2. Mitigating Manipulation and Enhancing Persuasion: A Reflective Multi-Agent Approach for Legal Argument Generation, [ResearchTrend.AI](#)
3. Enhancing Legal Argument Generation: From Basic Schemes to Trustworthy Multi-Agent Systems, [ISP Forum](#), University of Pittsburgh

SKILLS & QUALIFICATIONS

- **Research:** Statistical learning, LLM evaluation and benchmarking, multi-agent systems, interpretability, and document intelligence
- **Programming:** Python (PyTorch, Scikit-learn, Pandas and NumPy), R, C/C++, MATLAB, SQL, JavaScript
- **Cloud Technologies:** SageMaker, DynamoDB and API Gateway
- **Finance & Law:** PRC Legal Professional Qualification; two years of cross-border financing and credit transaction experience in commercial banking
- **Others:** L^AT_EX, Markdown, Linux, Shell and Anaconda