

Li Zhang

+1 412-328-6886 | liz239@pitt.edu

EDUCATION

University of Pittsburgh

Doctor of Philosophy in Intelligent Systems

Pittsburgh, PA, U.S.A.

Aug. 2024 - Apr. 2028

- Research: AI & Law
- Advisor: Prof. [Kevin Ashley](#)
- Courses: Large Language Models (Carnegie Mellon University), Natural Language Processing

Tsinghua University

Master of Computational Law

Beijing, China

Sep. 2019 - Jun. 2022

- Research: Computational Law

Australian National University

Study Abroad in College of Engineering and Computer Science

Canberra, Australia

Jul. 2017 - Dec. 2017

- GPA: 6.5/7.0

China University of Mining and Technology

Bachelor of Electrical Engineering

Jiangsu, China

Sep. 2015 - Jun. 2019

- GPA: 90.91/100, Ranking: 2/278

TECHNICAL EXPERIENCE

Thomson Reuters Labs

Applied Research Intern

Eagan, MN, U.S.A.

May 2026 - Aug. 2026

- Developed long-term memory capabilities for AI agents in legal-tech applications, supporting deep research over legal documents and long-form content generation for litigation and transactional workflows.
- Built and evaluated Legal Agent pipelines across dataset optimization, model adaptation, and automated assessment, working with science, product, and engineering partners to improve grounded legal workflows.

LEGAL EXPERIENCE

The Export-Import Bank of China, Jiangsu Branch

Legal Counselor

Nanjing, China

Jun. 2022 - Jul. 2024

- Led legal negotiations for complex banking transactions, focusing on regulatory compliance, risk mitigation, and cross-border financing matters in a highly regulated commercial banking environment.
- Drafted and reviewed credit and security documents while advising internal stakeholders on operational risk, transaction structuring, and dispute resolution strategy for banking operations.

SELECTED PUBLICATIONS

1. DLawBench: Evaluating LLMs Through Multi-Turn Legal Consultation [\[Link\]](#)

Li Zhang, Yuzhen Shi, Yiran Hu, Jingwen Zhang, Wenbo Lv, Yubo Ma, Wei Wang, Rongyao Shi, Yuanyang Qiu, Xinran Xu, Yuemeng Qi, Linlin Miao, Jaromir Savelka, Yun Liu, Kevin Ashley, Bing Zhao, Hu Wei, Lin Qu
EMNLP under review

2. Thinking Longer, Not Always Smarter: Evaluating LLM Capabilities in Hierarchical Legal Reasoning [\[Link\]](#)

Li Zhang, Matthias Grabmair, Morgan Gray, Kevin Ashley
CS&Law 2026

3. PLawBench: A Rubric-Based Benchmark for Evaluating LLMs in Real-World Legal Practice [\[Link\]](#)

Yuzhen Shi, Huanghai Liu, Yiran Hu, Song gaojie, Xu Xinran, Yubo Ma, Tianyi Tang, **Li Zhang**, Qingjing Chen, Feng Di, Wenbo Lv, Weiheng Wu, Kexin Yang, Sen Yang, Wei Wang, Rongyao Shi, Qiuyuan Yang, Yuemeng Qi, Jingwen Zhang, Xiaoyu Sui, Yifan Chen, Yi Zhang, An Yang, Bowen Yu, Dayiheng Liu, Junyang Lin, Weixin Shen, Bing Zhao, Charles Clarke, Hu Wei
ACL 2026

4. **DenseMixer: Improving MoE Post-Training with Precise Router Gradient** [\[Link\]](#)
Feng Yao, Junxia Cui, Ruohan Zhang, Liyuan Liu, Shibo Hao, **Li Zhang**, Chengyu Dong, Shuohang Wang, Yelong Shen, Jianfeng Gao, Jingbo Shang
NeurIPS 2025 Workshop on Structured Probabilistic Inference & Generative Modeling
5. **Do LLMs Truly Understand When a Precedent Is Overruled?** [\[Link\]](#)
Li Zhang, Jaromir Savelka, Kevin Ashley
JURIX 2025, Best Student Paper
6. **Mitigating Manipulation and Enhancing Persuasion: A Reflective Multi-Agent Approach for Legal Argument Generation** [\[Link\]](#)
Li Zhang, Kevin Ashley
ICAIL 2025 Workshop
7. **Measuring Faithfulness and Abstention: An Automated Pipeline for Evaluating LLM-Generated 3-ply Case-Based Legal Arguments** [\[Link\]](#)
Li Zhang, Morgan Gray, Jaromir Savelka, Kevin Ashley
ICAIL 2025 Workshop
8. **Generating Case-Based Legal Arguments with LLMs** [\[Link\]](#)
Morgan Gray, **Li Zhang**, Kevin Ashley
CS&Law 2025

PROJECTS

- | | |
|---|----------------------------|
| Benchmarking LLM Long-Context Understanding and Temporal Logic | Pittsburgh, PA, U.S.A. |
| <i>Li Zhang, Jaromir Savelka, Kevin Ashley</i> | <i>Sep. 2024 – Present</i> |
| <ul style="list-style-type: none">• Constructed a high-stakes benchmark dataset from real-world legal corpora to rigorously evaluate LLMs' ability to comprehend long-context documents and track complex overruling relationships.• Analyzed temporal biases and hallucinations in model outputs, developing specific metrics to ensure AI-generated case histories accurately reflect the temporal dynamics of legal precedents. | |
| Reflective Multi-Agent Frameworks for Ethical Legal Persuasion | Pittsburgh, PA, U.S.A. |
| <i>Li Zhang, Kevin Ashley</i> | <i>Aug. 2024 – Present</i> |
| <ul style="list-style-type: none">• Architected a multi-agent AI system employing reflective verification processes, where separate agents critique and verify generated arguments against factual records to minimize hallucination.• Implemented ethical guardrails to ensure AI assistants advocate effectively without manipulating facts or fabricating evidence, aligning system outputs with professional legal standards. | |
| Evaluating and Enhancing Hierarchical Legal Reasoning | Pittsburgh, PA, U.S.A. |
| <i>Li Zhang, Matthias Grabmair, Morgan Gray, Kevin Ashley</i> | <i>Nov. 2024 – Present</i> |
| <ul style="list-style-type: none">• Investigated the phenomenon of reasoning collapse in LLMs, specifically evaluating their performance in multi-stage legal reasoning tasks that require distinguishing cases based on subtle factual predicates.• Developed hierarchical inference strategies to enhance the robustness of legal AI, improving the model's ability to perform nuanced logical deductions in complex case law scenarios. | |

TALKS

1. Do LLMs Truly Understand When a Precedent Is Overruled, [2025 Pittsburgh Legal AI Summit](#), Carnegie Mellon University
2. Mitigating Manipulation and Enhancing Persuasion: A Reflective Multi-Agent Approach for Legal Argument Generation, [ResearchTrend.AI](#)
3. Enhancing Legal Argument Generation: From Basic Schemes to Trustworthy Multi-Agent Systems, [ISP Forum](#), University of Pittsburgh

SKILLS & QUALIFICATIONS

- **Programming:** Python (PyTorch, Scikit-learn, Pandas and NumPy), SQL, JavaScript
- **Cloud Technologies:** SageMaker, DynamoDB and API Gateway
- **Legal:** PRC Legal Professional Qualification
- **Others:** $\text{\LaTeX}$, Markdown, Linux, Shell and Anaconda